

PLAN FIRST, ACT LATER: STREAMING REASONING FOR EXECUTABLE MEMORY IN ROBOT CONTROL

Anonymous authors

Paper under double-blind review

ABSTRACT

Memory-dependent manipulation requires a robot to retain task-relevant evidence beyond its short observation window. Existing memory-augmented policies typically retrieve relevant information from history and reason over it at execution time, placing a heavy reasoning burden on the high-level model. To address this challenge, we propose **PAL** (Plan first, Act Later), a hierarchical memory system designed to alleviate this reasoning burden. It maintains a short-term visual buffer, a semantic-event memory, and an executable `subtask_list`. A high-level vision-language model (VLM) converts streaming visual observations into the `subtask_list`. At execution time, the VLM selects the currently relevant subtask from the `subtask_list` and passes it to a low-level vision-language-action (VLA) policy. To teach the model what to write to memory and when to update it, we introduce a two-stage procedure that first densifies sparse update supervision and then uses self-labeling with task-grounded verification to generate training data that better match the state distribution encountered during execution. We validate the proposed method on memory-dependent benchmarks and real-world manipulation tasks. On RoboMME, **PAL** achieves a **56.05%** success rate, compared with 44.51% for the strongest baseline, FrameSamp-Modul. In real-world experiments across four manipulation tasks, **PAL** reaches **60.83%** overall success, outperforming MemER (45.00%) and FrameSamp-Modul (28.33%). These results demonstrate that proactively writing future-relevant subtasks into hierarchical memory improves long-horizon visual reasoning and robot control.




Figure 1: Overview of our pipeline. The high-level VLM incrementally converts streaming visual observations into hierarchical memory and retrieves the currently relevant subtask for execution by the low-level VLA policy.

1 INTRODUCTION

Vision-language-action (VLA) models have made substantial progress toward general-purpose robot control (Intelligence et al., 2025b; Kim et al., 2024). By grounding language instructions in visual observations and actions, VLA provides a flexible interface for solving diverse manipulation tasks in generalizable open-world. However, many existing policies still rely on a Markovian or near-Markovian view of the world: the action to execute is assumed to be recoverable from the current observation or a short window of recent frames. This assumption becomes fragile in memory-dependent settings, where successful behavior may depend on historical information observed much earlier. For example, repetitive manipulation requires remembering progress (Yang et al., 2026; Dai et al., 2026; Chen et al., 2026a) and active exploration requires retrieving the location or state of an object discovered earlier when the robot revisits the scene (Shah et al., 2026; Sridhar et al., 2025).

To address these memory-dependent challenges, recent works augment VLA models with a memory bank that retains task-relevant information from past observations. At execution time, the model retrieves relevant entries from the bank to support stable long-term decision making (Torne et al., 2026; Sridhar et al., 2025; Shi et al., 2026). Although this avoids repeatedly processing the raw history, the high-level model must still search across a large memory bank and infer the next subtask at inference time. Such banks often contain redundant information, making it difficult for the model to fetch the required information and lead to more task failures.

One key observation is that: ***Evidence for defining a future goal are available before executing the goal.*** As illustrated in Fig. 1, the task requires the robot to press the button and then pick up the highlighted cube. While the robot is still pressing the button, the model can store the future subtask “pick up the cube” in the subtask list when the cube is highlighted. When the transient highlight disappears, the robot can simply retrieve the stored subtask for execution, without re-identifying which cube to pick up. This pattern occurs in many memory-dependent robotic tasks, where the future plan can be determined once the necessary evidence becomes available.

Motivated by this view, we propose **Plan First, Act Later (PAL)**, a hierarchical memory pipeline for VLA long-term decision-making. It comprises a three-level memory system: a short-term visual buffer, a semantic-event memory, and an executable `subtask_list`. These three memory components have distinct roles: the visual buffer retains recent observations, the semantic-event memory maintains task-relevant events, and the executable `subtask_list` stores future-relevant subtasks text instructions. At each query, a high-level vision-language model (VLM) updates the event memory and the `subtask_list`, then selects the subtask to execute for the low-level VLA. This design has two benefits: (1) first, streaming memory updates incrementally compress the visual history, reducing the context required for long-horizon reasoning; (2) second, the stored subtasks can be directly retrieved without searching across a large memory bank for subtask inference.

To teach the model what to write to memory and when to update it, we propose a two-stage training pipeline. The pipeline combines dense update-timing supervision with self-labeling and a task-grounded verifier, allowing the training data to better reflect the model’s execution rollout distribution. Evaluation results on RoboMME (Dai et al., 2026), a benchmark of long-horizon, memory-dependent manipulation tasks, show that **PAL** achieves a **56.05%** success rate, compared with 44.51% for the strongest baseline, FrameSamp-Modul. In real-world manipulation experiments across four tasks, **PAL** reaches **60.83%** overall success, outperforming MemER (45.00%) and FrameSamp-Modul (28.33%), demonstrating the practical benefits of our approach. Our contributions are as follows:

- We introduce **PAL**, a hierarchical memory architecture with a streaming-style reasoning mechanism for VLA long-term decision-making. It progressively transforms streaming visual observations into semantic events and executable subtasks, providing a structured memory for diverse long-horizon, memory-dependent tasks.
- We develop a two-stage training pipeline that enables the model to learn reliable hierarchical-memory update behavior. It produces accurate and timely memory updates, yielding reliable executable subtask lists for long-horizon reasoning.
- We demonstrate the effectiveness of **PAL** through evaluations on long-horizon, memory-dependent manipulation tasks, where it delivers consistent gains over competitive VLA and memory-augmented methods in both simulated and real-world settings.

2 RELATED WORKS

Vision-language-action models for robot control. Recent VLA models have evolved from instruction-conditioned visuomotor policies into general robot control systems. Early efforts such as OpenVLA and $\pi_{0.5}$ established a unified interface that maps language instructions and visual observations to low-level actions (Kim et al., 2024; Intelligence et al., 2025b). More recent systems improve generalization through larger-scale training, broader robot data, and continued experience-driven learning (Intelligence et al., 2025a; 2026; Team et al., 2026). Despite this progress, most VLA policies still operate on the current observation or a short recent context, motivating hierarchical planning and memory mechanisms for tasks that require information from earlier history.

Hierarchical planning and subgoal-conditioned control. Long-horizon robot tasks are often handled with a hierarchical interface, where a high-level model selects a subtask or subgoal for a low-level policy to execute (Intelligence et al., 2025b; Shi et al., 2025; Chen et al., 2026a). This interface separates selecting the next subtask from generating the low-level actions that execute it. τ_0 -VLA extends this hierarchy with execution memory and world-model-guided test-time computation, predicting and comparing candidate subtask outcomes before commitment (Cai et al., 2026). Our work adopts the similar interface: the low-level policy executes the current subtask, and the high-level model maintains the semantic state needed to select future subtasks. However, the hierarchical interface alone does not specify how to use information from earlier observations, which motivates memory-augmented VLA methods.

Memory-augmented VLA models and benchmarks. Memory-augmented VLA methods address this missing component by storing and retrieving historical information for high-level decision making. Recent works use amortized multi-frame context (Jang et al., 2025), multi-scale embodied memory (Torne et al., 2026), experience retrieval (Sridhar et al., 2025), keyframe-based memory (Chen et al., 2026b), structured scene memory (Ju et al., 2026), and perceptual-cognitive token memories (Shi et al., 2026). HiMe introduces hierarchical embodied memory for long-horizon VLA control, providing another memory-bank approach (Ji et al., 2026). In parallel, benchmarks such as RoboMME, RMBench, and RoboMemArena highlight that robotic policies often fail when task success requires retaining and using historical information (Dai et al., 2026; Chen et al., 2026a; Lei et al., 2026). Although these methods make past information accessible, they still require the high-level model to search a large memory bank at execution time, so the computational cost grows with the amount of stored history. In contrast, our method allows the model to record a future-relevant subtask as soon as the supporting evidence is observed, before the subtask must be executed, and later select it from the editable `subtask_list`. Concurrent work, MaP-WAM (Zhao et al., 2026), similarly plans ahead and stores future-relevant information in memory. However, MaP-WAM plans only the next step, whereas our editable `subtask_list` retains all inferred future subtasks and serves as a persistent component of the memory system.

3 PRELIMINARIES

Problem setup. We consider memory-dependent VLA tasks in which the correct behavior at timestep t may depend on information that is not visible in the current observation. Let L denote the language instruction, o_t the visual observation, a_t the low-level action, and $H_t = \{(o_i, a_i)\}_{i=1}^t$ the raw interaction history. Long-horizon tasks are increasingly challenging for a fully history-based policy $\pi(a_t | H_t, L)$ since the visual context grows with time and contains many redundant frames.

Hierarchical memory formulation. We organize task memory at three temporal scales. The short-term visual buffer $\mathcal{V}_t$ stores recent RGB images from the current observation window. The other two components are ordered lists of natural-language entries. The semantic-event memory $\mathcal{E}_t$ records events that have occurred in the environment and may be relevant to the task, with one event per entry. The executable memory $\mathcal{P}_t$ is the `subtask_list`, and each of its entries is one subtask. An example is shown in Fig. 1. Together, these components form the hierarchical memory state

$$\mathcal{M}_t := (\mathcal{V}_t, \mathcal{E}_t, \mathcal{P}_t),$$

At each query, the VLM uses the incoming visual observations and the hierarchical memory to update either list and to select the currently subtask ℓ_t , which is passed to the low-level policy:

$$a_t = \pi_{\text{VLA}}(o_t, \ell_t, L).$$

Plan first, Act later. Our approach to constructing the subtask list is inspired by a simple temporal separation: when visual evidence reveals that a subtask will be needed before it is time to execute it, the model can store the subtask for later execution. Formally, a subtask written to memory at t_1 can be selected as ℓ_{t_2} for execution at a later time t_2 :

$$\mathcal{M}_{t_1+\Delta t} = g_\theta(\mathcal{M}_{t_1}, o_{t_1:t_1+\Delta t}, L), \quad \ell_{t_2} \in \mathcal{P}_{t_1+\Delta t}, \quad t_1 \leq t_2.$$

4 METHODOLOGY

4.1 OVERVIEW

Our methodology has two parts. Section 4.2 describes the memory construction and querying mechanism used whenever the VLM runs, including both training rollouts and inference. Section 4.3 then describes how we train the VLM to acquire this capability: learning what to write to memory, when to update it, and which stored subtask to execute.

4.2 STREAM-BASED MEMORY REASONING

Following the three-level hierarchical memory formulation introduced in Sec. 3, we propose a streaming-style reasoning pipeline that processes history as a stream of image chunks, as illustrated in Fig. 2. For a query at time t , let $\mathcal{C}_t = o_{t-\Delta t:t}$ denote the observations collected since the previous query. The VLM takes the previous hierarchical memory $\mathcal{M}_{t-\Delta t}$, the current observation chunk $\mathcal{C}_t$, and the task instruction L as input. The visual buffer stored in $\mathcal{M}_{t-\Delta t}$ is combined with $\mathcal{C}_t$ and downsampled to a fixed number of frames before being passed to the VLM. The VLM then outputs three fields. The quantity $\Delta\mathcal{E}_t$ is a newly predicted event, while $\Delta\mathcal{P}_t$ is a newly predicted subtask to be appended to the subtask list. The third output, ℓ_t , is the current subtask passed to the low-level VLA. Because image chunks arrive incrementally, the evidence available at a query may be insufficient to determine whether a new event has occurred or which future subtask should be appended. In this case, the VLM outputs an empty $\Delta\mathcal{E}_t$ or $\Delta\mathcal{P}_t$, respectively, and the corresponding memory is retained until further observations become available. The memory transition is written as

$$(\Delta\mathcal{E}_t, \Delta\mathcal{P}_t, \ell_t) = g_\theta(\mathcal{M}_{t-\Delta t}, \mathcal{C}_t, L),$$

For ordered sequences, $\text{Append}(A, B)$ denotes appending B to the end of A . The subtask list is updated as

$$\mathcal{P}_t = \text{Append}(\mathcal{P}_{t-\Delta t}, \Delta\mathcal{P}_t).$$

If $\Delta\mathcal{P}_t$ is non-empty, the previous event memory $\mathcal{E}_{t-\Delta t}$ is cleared; otherwise, it is retained. Let $\tilde{\mathcal{E}}_{t-\Delta t}$ denote this retained or reset event memory:

$$\tilde{\mathcal{E}}_{t-\Delta t} = \begin{cases} \emptyset, & \text{if } \Delta\mathcal{P}_t \neq \emptyset, \\ \mathcal{E}_{t-\Delta t}, & \text{otherwise.} \end{cases}$$

The event memory is updated by appending the predicted event:

$$\mathcal{E}_t = \text{Append}(\tilde{\mathcal{E}}_{t-\Delta t}, \Delta\mathcal{E}_t).$$




Figure 2: Streaming memory update.




Figure 3: Dense update-timing supervision. A sparse update annotation is expanded into multiple training samples with visual buffers of different lengths, exposing the model to different levels of evidence around the reference update boundary.

The visual buffer is cleared after a memory update; otherwise, the new chunk is appended to the previous visual buffer:

$$\mathcal{V}_t = \begin{cases} \emptyset, & \text{if } \Delta \mathcal{E}_t \neq \emptyset \text{ or } \Delta \mathcal{P}_t \neq \emptyset, \\ \text{Append}(\mathcal{V}_{t-\Delta t}, C_t), & \text{otherwise.} \end{cases}$$

Clearing the buffer after a memory update prevents visual evidence from accumulating excessively and making subsequent reasoning harder. After the subtask list is updated, the event buffer is also cleared so that subsequent event recognition starts relative to the updated list. The complete update procedure for one query is summarized in Algorithm 1 in the appendix.

4.3 TWO-STAGE TRAINING WITH SELF-LABELING AND VERIFICATION

We propose a two-stage training procedure to teach the model what to write to the subtask list and when to update it, while jointly training it to predict the current subtask at each query. Stage 1 augments the training data by creating additional training examples from each sparse annotation. However, this augmentation changes the training data distribution, introducing a mismatch with the model’s inference distribution. We therefore introduce Stage 2 to regularize the model’s memory-update behavior and reduce this mismatch.

Stage 1: Dense update-timing supervision. As illustrated in Figure 3, we first construct a dataset of sparse annotations; details of the simulation and real-world dataset construction are provided in Appendices C.2 and D.1, respectively. Each sparse annotation is provided at a single timestep t_i . We denote the training labels for the new event and subtask at this timestep by $\Delta \mathcal{E}_i^{\text{ref}}$ and $\Delta \mathcal{P}_i^{\text{ref}}$, respectively. Such sparse annotation is data-inefficient because it provides only one supervision point for each update. Moreover, visual evidence changes only slightly between neighboring queries, making the update and no-update cases difficult for the model to distinguish.

To address this issue, we sample multiple query timesteps t at regular intervals from $[t_{i-1}, \frac{t_i + t_{i+1}}{2})$. For each query, we construct $\mathcal{V}_t$ from the observations accumulated since the preceding update up to timestep t , so later queries contain more visual evidence. For each sampled query timestep, we assign $\emptyset$ to the event and subtask-list targets when $t < \frac{t_{i-1} + t_i}{2}$, and assign $\Delta \mathcal{E}_i^{\text{ref}}$ and $\Delta \mathcal{P}_i^{\text{ref}}$ when $t \geq \frac{t_{i-1} + t_i}{2}$. We use this simplified assignment as a practical rule: queries presumed to have insufficient information for an update receive empty targets, whereas later queries receive the corresponding reference event and subtask labels. The effectiveness of this simple approximation is further examined in Sec. 5.2. In this way, the single annotation at t_i is expanded into annotations over a range of query timesteps. The model is fine-tuned on these densified labels to learn the memory update schema and approximate update timing from the reference teacher’s distribution.

Stage 2: Self-labeling with task-grounded verification. Despite this supervision, the Stage 1 model may still produce updates too frequently or at imprecise times. In particular, it may update before enough evidence has accumulated or repeat content that is already in memory. Stage 2 is designed to address these issues through task-grounded verification.

Concretely, we roll out the Stage 1 model on the training set, maintaining the memory states generated by the model. Let $\frac{t_{i-1}+t_i}{2} \leq t < \frac{t_i+t_{i+1}}{2}$ be the timing band of the i -th reference update, whose reference event and subtask-list increments are $\Delta\mathcal{E}_i^{\text{ref}}$ and $\Delta\mathcal{P}_i^{\text{ref}}$. Let $\mathcal{X} \in \{\mathcal{E}, \mathcal{P}\}$ stand for either memory. For each prediction, the verifier checks the corresponding reference label and whether that memory already stores it:

$$\Delta\mathcal{X}_t \leftarrow \emptyset, \quad \text{if } \Delta\mathcal{X}_t \neq \Delta\mathcal{X}_i^{\text{ref}} \text{ or } \Delta\mathcal{X}_t \in \mathcal{X}_{t-\Delta t}.$$

Thus, an incorrect prediction, as well as one that duplicates an entry already stored in memory, is replaced with $\emptyset$, while the other prediction is unaffected. This output is then used to update the memory following the update mechanism in Section 4.2, and the updated memory state $\mathcal{M}_t$ is the input of the next query. If

$$t \geq \frac{t_i + t_{i+1}}{2} \quad \text{and} \quad \Delta\mathcal{X}_i^{\text{ref}} \notin \mathcal{X}_t \text{ for some } \mathcal{X} \in \{\mathcal{E}, \mathcal{P}\},$$

the rollout is marked as failed, since the model has not produced and stored the i -th reference event or subtask-list update by the deadline. We retain only rollouts in which the verifier accepts an update at every required update point. The resulting update annotations remain sparse, and their timing is determined within a limited range by the VLM. As a result, the data distribution is closer to the model’s output distribution and is easier for the model to learn. The model is then fine-tuned on these self-labeled trajectories. The complete procedure is given in Algorithm 2 in the appendix.

5 EXPERIMENTS

This section answers the following questions:

- By how much does **PAL** improve success rates on memory-dependent manipulation tasks compared with baselines?
- Why does **PAL** perform better than baselines on long-horizon tasks?
- How do the two training stages contribute to **PAL**’s overall performance?
- How does **PAL** perform in real-world robotic manipulation tasks compared with baselines?

5.1 SIMULATED EXPERIMENTS

RoboMME benchmark and baselines. We evaluate our model on the RoboMME benchmark (Dai et al., 2026), which is a collection of long-horizon, memory-dependent tasks for robotic manipulation. We compare **PAL** with five baselines. $\pi_{0.5}$ uses only the current observation and task instruction, providing a memory-free reference without access to observation history. **FrameSamp-Modul** (Dai et al., 2026) uniformly samples frames from the observation history and injects their visual features into the policy through feature modulation, providing a perceptual-memory baseline without explicit high-level reasoning. **MemER** (Sridhar et al., 2025) retrieves relevant experience from the observation history and uses a VLM to produce a grounded subtask for the low-level VLA. **PAL** adopts the same high-level–low-level control interface, and the low-level VLA used by MemER is identical to that of **PAL**. **EventVLA** (Yang et al., 2026) is an end-to-end VLA that learns to retain sparse, task-relevant keyframes in its Keyframe Evidence Memory. It provides a complementary comparison between **PAL**’s semantic, executable memory and learned visual evidence memory. Finally, **GPT-6-astra + VLA** replaces **PAL**’s fine-tuned high-level VLM with GPT-6-astra while using the same low-level VLA, and represents its memory as a language description of the current environment from which it predicts the next subtask. Table 1 reports group-level means. The full per-task breakdown and the implementation and baseline details are provided in Appendices C.1 and C.3, respectively.

Table 1: Mean success rates (%) on RoboMME task groups. **Short-video** and **Long-video** infer an executable subtask list from conditioning videos of different lengths; **In-context** adapts it from observations during the episode; and **Episode-start** has all information needed to infer the complete list at the beginning of the episode. The AVG column reports mean $\pm$ standard deviation when multiple evaluation seeds are available.

Method	Short-video	Long-video	In-context	Episode-start	AVG
Baselines					
$\pi_{0.5}$	10.67	28.45	11.72	36.15	17.93 ± 0.61
EventVLA	27.60	19.00	32.50	21.33	26.57 ± 0.13
FrameSamp-Modul	41.87	46.00	27.05	72.96	44.51 ± 0.77
MemER	37.52	28.00	41.00	65.11	42.38 ± 0.33
GPT-6-astra + VLA	31.71	66.00	28.00	56.67	39.75
Streaming Reasoning					
Stage 1 Only	31.43	17.40	32.50	70.94	37.35 ± 1.77
Sparse Label	36.74	30.44	38.64	73.87	43.39 ± 2.26
PAL (ours)	45.26	81.60	46.20	77.33	56.05 ± 1.38

Results and analysis. The main results are summarized in Table 1, with complete per-task results provided in Appendix C.1. We analyze these results through the following research questions.

Q1: By how much does PAL improve success rates on memory-dependent manipulation tasks compared with baselines?

PAL achieves an overall success rate of 56.05%, improving upon the strongest baseline, FrameSamp-Modul (44.51%), by 11.54 percentage points. We focus the comparison on MemER and GPT-6-astra + VLA because, like **PAL**, both use a hierarchical control structure in which a high-level model generates a subtask for a low-level VLA. **PAL** outperforms MemER (42.38%) and GPT-6-astra + VLA (39.75%) by 13.67 and 16.30 percentage points, respectively, and achieves the highest success rate on 4 of the 16 RoboMME tasks.

Q2: Why does PAL perform better than baselines on long-horizon tasks?

We attribute **PAL**’s gains to two complementary components: the executable subtask list decouples subtask inference from execution-time activation, while streaming memory updates reduce the visual context and VLM capacity required for long-horizon reasoning.

The **Episode-start** group directly evaluates the *benefit of constructing* the subtask list ahead of execution, as all information needed to infer the complete list is available at the beginning of the episode. **PAL** achieves 77.33% on this group, outperforming the strongest baseline, FrameSamp-Modul (72.96%), by 4.37 percentage points. For a more direct comparison, MemER and GPT-6-astra + VLA also use a high-level model to generate a subtask for the low-level VLA. However, **PAL** constructs the complete subtask list in advance, whereas MemER generates a subtask only when it is needed. **PAL** outperforms MemER (65.11%) by 12.22 percentage points, indicating the benefit of planning future subtasks ahead of execution rather than inferring them on demand.

The remaining groups further demonstrate the *benefit of streaming memory* updates. **PAL** achieves success rates of 45.26%, 81.60%, and 46.20% on these groups, improving over FrameSamp-Modul (41.87%, 46.00%, and 27.05%) by 3.39, 35.60, and 19.15 percentage points, respectively. The largest gain is on **Long-video**, where the conditioning stream is longest. **PAL**’s streaming updates compress this visual history into a compact subtask list, so the VLM reasons over a smaller context rather than the full video.

5.2 ABLATION STUDY

To validate the effectiveness of our two-stage training approach, we conduct ablation experiments on the RoboMME benchmark. We compare four variants: (1) **Sparse Label**, trained on data that annotate event or subtask-list updates only at the end of each evidence chunk; (2) **Stage 1 Only**, which uses only the densified supervision from the reference teacher without self-labeling; (3) **Update Gate**, which uses the Stage 1 checkpoint together with an additional ground-truth update gate

Table 2: Subtask-list update behavior. For each task, we report the average number of `subtask_list` updates per episode (Update freq.) and the accuracy of the final list (Final acc.). Parentheses report the relative change in update frequency from ground truth.

Variant	PatternLock		RouteStick		VideoPlaceOrder		VideoPlaceButton	
	Update freq.	Final acc.	Update freq.	Final acc.	Update freq.	Final acc.	Update freq.	Final acc.
Sparse Label	2.68 (-14.6%)	22.00%	3.34 (-8.7%)	50.00%	0.00 (-100.0%)	0.00%	0.00 (-100.0%)	0.00%
Stage 1 Only	4.32 (+37.6%)	24.00%	5.28 (+44.3%)	30.00%	1.34 (+34.0%)	18.00%	1.38 (+38.0%)	18.00%
Update Gate	3.10 (-1.3%)	78.00%	3.66 (0.0%)	82.00%	1.00 (+0.0%)	80.00%	1.00 (+0.0%)	78.00%
Ground truth	3.14 (0.0%)	—	3.66 (0.0%)	—	1.00 (0.0%)	—	1.00 (0.0%)	—
PAL	2.94 (-6.4%)	54.00%	3.46 (-5.5%)	68.00%	1.04 (+4.0%)	84.00%	1.02 (+2.0%)	80.00%

provided by the environment at inference time. Similar to the gate in Sec. 4.3, it opens an update slot when each interval reaches 80% of its duration; the slot remains open until the VLM produces its first non-empty update, and any update outside the active slot is discarded. (4) **PAL** is our complete approach with both stages and streaming updates. All of them use the same hierarchical memory architecture and low-level VLA policy as **PAL**. The only difference is the training procedure and the update gate used at inference time.

Q3: How do the two training stages contribute to **PAL’s overall performance?** Table 1 reports group-level averages, with the full per-task results in Appendix C.1. Stage 1 Only achieves 37.35% overall success, Sparse Label improves to 43.39%, and **PAL** reaches 56.05%. To explain these differences, we examine the update behavior in Table 2 on four representative streaming-reasoning tasks. PatternLock and RouteStick involve a large number of list updates as evidence arrives, whereas VideoPlaceOrder and VideoPlaceButton require evidence to be accumulated over a long video stream. Comparing Stage 1 Only with Update Gate, we find that applying the update gate greatly improves final-list accuracy. This is because Update Gate provides the model with approximate update timings, thereby decoupling when to update from what to update. The improvements on the final-list accuracy implies that Stage 1 helps the model learn what content to update, but not when to update it.

We then calculate the update frequency of each variant. As expected, Stage 1 Only tends to update the executable list too often, since the model is trained with a dataset that over densifies the update timing. Sparse Label under-updates and even fails to update any subtask list for Long-video tasks, which implies sparse supervision is insufficient for learning the correct update behavior. **PAL**’s update frequency is closer to ground truth and yields higher final-list accuracy. This suggests that Stage 2 regularizes the overly frequent update behavior learned in Stage 1, producing more accurate memory states for long-horizon reasoning.

5.3 REAL-WORLD EXPERIMENTS

We evaluate our pipeline in the real world using a Franka Research 3 robot arm. The evaluation comprises four tasks adapted from their simulated counterparts, each targeting a distinct memory capability. PickXtimes requires the robot to pick a green box a specified number of times, testing its ability to track task progress. In VideoUnmask, the model observes a video showing the layout of colored cubes and must subsequently point to the requested cube. VideoPlaceOrder tests long-video reasoning by requiring the model to recall where a target box was placed earlier in the video. PatternLock evaluates imitation and reasoning over a temporally dense action sequence. For each task, we evaluate 10 episodes spanning different difficulty levels and repeat each episode three times, for a total of 30 evaluations per task. We compare against the MemER and FrameSamp-Modul baselines. Appendix D provides further details on the tasks, evaluation protocol, and rollout visualizations.




Figure 4: Real-world manipulation tasks used to evaluate **PAL**. The tasks require retaining visual evidence from previous observations to support subsequent manipulation.

Table 3: Real-world success rates (% , mean $\pm$ sample standard deviation over three repeats of each 10-episode evaluation).

Method	PatternLock	PickXtimes	VideoPlaceOrder	VideoUnmask	Overall
FrameSamp-Modul	30.00 $\pm$ 10.00	40.00 $\pm$ 30.00	13.33 $\pm$ 5.77	30.00 $\pm$ 0.00	28.33 $\pm$ 3.82
MemER	33.33 $\pm$ 5.77	0.00 $\pm$ 0.00	46.67 $\pm$ 30.55	100.00 $\pm$ 0.00	45.00 $\pm$ 6.61
Ours	40.00 $\pm$ 0.00	53.33 $\pm$ 15.28	50.00 $\pm$ 0.00	100.00 $\pm$ 0.00	60.83 $\pm$ 3.82

Q4: How does **PAL perform in real-world robotic manipulation tasks compared with baselines?** The results are reported in Table 3. FrameSamp-Modul remains weak across all four tasks and almost entirely fails on VideoPlaceOrder (13.33%), indicating that a short, evenly sampled visual window is insufficient for recalling occluded events in long videos. MemER performs well on VideoUnmask (100.00%), a comparatively short-context recall task, but is weaker than our method on longer-context tasks such as PickXtimes (0.00% vs. 53.33%) and VideoPlaceOrder (46.67% vs. 50.00%), and remains limited on PatternLock (33.33%). In contrast, our method attains the highest overall success rate of 60.83%, with the largest gains on the longer-horizon tasks. These results suggest that converting streaming observations into hierarchical, executable memory preserves the task-relevant events needed for real-world long-horizon manipulation, rather than relying on lossy frame sampling or retrieval from raw history.

6 CONCLUSION

We presented **PAL**, a hierarchical memory pipeline that incrementally converts streaming visual observations into semantic events and a subtask list before they are needed for control. By separating early subtask construction from execution-time selection, **PAL** reduces the need to repeatedly revisit a large visual history. We further introduced a two-stage training procedure that combines dense update supervision with verified self-labeling to learn reliable memory updates under the model’s rollout distribution. Experiments on RoboMME and four real-world manipulation tasks demonstrate that **PAL** improves memory-dependent robot control, particularly when relevant evidence must be retained over long horizons.

While **PAL** improves upon prior work, it still has limitations. Its executable memory is represented as a sequential subtask list and therefore cannot encode conditional branches during execution. Programmatic planning methods such as Code as Policies demonstrate that executable programs can express conditional control flow and perception-dependent behavior (Liang et al., 2023). Future work could extend **PAL** with branchable executable memory that selects different subtask sequences according to execution-time observations.

AI USE STATEMENT

In this work, we used generative AI tools to generate synthetic data sets, design or provide feedback on research methodology or experiments, implement methods, assist with translation, clean and reformat datasets, support qualitative and thematic data analysis, and interpret results. We have

not used generative AI tools to help develop theoretical models or conceptual pipelines, formulate mathematical claims, or propose or refine hypotheses. Providing critical ingredients for proving mathematical claims and assisting in the writing of proofs are not applicable to this work. Additionally, we used generative AI tools to create or modify scientific figures or images, draft parts of the research paper, summarize or analyse existing literature, brainstorm, source and search for information, edit the research paper to improve readability, identify relevant literature, format references, suggest a structure for the research paper, and propose a title or keywords for the research paper. We have reviewed all AI-assisted work: generated data and code were checked by the authors, and all drafted text was revised by the authors. We take responsibility for the final content of this work, including text, claims or artifacts produced with the aid of generative AI.

ETHICS STATEMENT

This work does not raise ethical concerns. It studies memory-dependent robot manipulation in simulation and on a robot arm in a controlled laboratory setting, involves no human subjects, and we do not foresee harmful applications.

REPRODUCIBILITY STATEMENT

All code and data needed to reproduce this work are available at <https://anonymous.4open.science/w/pal-webpage-2141/>. The streaming memory update and the Stage 2 self-labeling procedure are given as pseudocode in Appendix 1 and Appendix 2. The network architecture and training hyperparameters are listed in Appendix B. The construction of the simulated and real-world datasets is described in Appendices C.2 and D.1, and the evaluation protocols are described in Appendices C.3 and D.3.

REFERENCES

- Xiaowei Cai, Yunuo Cai, Bingao Chen, Jingxiao Chen, Zhi Chen, Siyuan Feng, Tengyu Hou, Jingshun Huang, Han Jiang, Runkun Ju, Dong Li, Mingxiang Li, Shaowei Li, Xinchun Li, Yifan Li, Yi Liu, Zhongyuan Liu, Jianlan Luo, Junwen Miao, Ruiqi Ni, Buqing Nie, Mingjie Pan, Xinlin Ren, Jianheng Song, Jiaxu Wang, Peiqi Wang, Sen Wang, Xiaoyan Wang, Dafeng Wei, Dongming Wu, Pengwei Xie, Pu Yang, Hangjian Ye, Xiangyu Yue, Jinyu Zhang, Qinglin Zhang, Xueyong Zhao, Pengfei Zhou, and Yue Zhou. τ_0 -vla: a hierarchical robot foundation model with world-model-guided test-time computation, 2026. URL <https://arxiv.org/abs/2608.16885>.
- Tianxing Chen, Yuran Wang, Mingleyang Li, Yan Qin, Hao Shi, Zixuan Li, Yifan Hu, Yingsheng Zhang, Kaixuan Wang, Yue Chen, Hongcheng Wang, Renjing Xu, Ruihai Wu, Yao Mu, Yaodong Yang, Hao Dong, and Ping Luo. Rmbench: Memory-dependent robotic manipulation benchmark with insights into policy design, 2026a. URL <https://arxiv.org/abs/2603.01229>.
- Yipeng Chen, Wentao Tan, Lei Zhu, Fengling Li, Jingjing Li, Guoli Yang, and Heng Tao Shen. Non-markovian long-horizon robot manipulation via keyframe chaining, 2026b. URL <https://arxiv.org/abs/2603.01465>.
- Yinpei Dai, Hongze Fu, Jayjun Lee, Yuejiang Liu, Haoran Zhang, Jianing Yang, Chelsea Finn, Nima Fazeli, and Joyce Chai. Robomme: Benchmarking and understanding memory for robotic generalist policies, 2026. URL <https://arxiv.org/abs/2603.04639>.
- Physical Intelligence, Ali Amin, Raichelle Aniceto, Ashwin Balakrishna, Kevin Black, Ken Conley, Grace Connors, James Darpinian, Karan Dhabalia, Jared DiCarlo, Danny Driess, Michael Equi, Adnan Esmail, Yunhao Fang, Chelsea Finn, Catherine Glossop, Thomas Godden, Ivan Goryachev, Lachy Groom, Hunter Hancock, Karol Hausman, Gashon Hussein, Brian Ichter, Szymon Jakubczak, Rowan Jen, Tim Jones, Ben Katz, Liyiming Ke, Chandra Kuchi, Marinda Lamb, Devin LeBlanc, Sergey Levine, Adrian Li-Bell, Yao Lu, Vishnu Mano, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Allen Z. Ren, Charvi Sharma, Lucy Xiaoyang Shi, Laura Smith, Jost Tobias Springenberg, Kyle Stachowicz, Will Stoeckle, Alex Swerdlow, James Tanner, Marcel Torne, Quan Vuong, Anna Walling, Haohuan Wang, Blake Williams, Sukwon Yoo, Lili

- 540 Yu, Ury Zhilinsky, and Zhiyuan Zhou. $\pi_{0.6}^*$: a vla that learns from experience, 2025a. URL
541 <https://arxiv.org/abs/2511.14759>.
542
- 543 Physical Intelligence, Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess,
544 Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Manuel Y. Galliker, Dibya Ghosh,
545 Lachy Groom, Karol Hausman, Brian Ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Devin
546 LeBlanc, Sergey Levine, Adrian Li-Bell, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Allen Z.
547 Ren, Lucy Xiaoyang Shi, Laura Smith, Jost Tobias Springenberg, Kyle Stachowicz, James Tanner,
548 Quan Vuong, Homer Walke, Anna Walling, Haohuan Wang, Lili Yu, and Ury Zhilinsky. $\pi_{0.5}$: a
549 vision-language-action model with open-world generalization, 2025b.
- 550 Physical Intelligence, Bo Ai, Ali Amin, Raichelle Aniceto, Ashwin Balakrishna, Greg Balke, Kevin
551 Black, George Bokinsky, Shihao Cao, Thomas Charbonnier, Vedant Choudhary, Foster Collins,
552 Ken Conley, Grace Connors, James Darpinian, Karan Dhabalia, Maitrayee Dhaka, Jared DiCarlo,
553 Danny Driess, Michael Equi, Adnan Esmail, Yunhao Fang, Chelsea Finn, Catherine Glossop,
554 Thomas Godden, Ivan Goryachev, Lachlan Groom, Haroun Habeeb, Hunter Hancock, Karol
555 Hausman, Gashon Hussein, Victor Hwang, Brian Ichter, Connor Jacobsen, Szymon Jakubczak,
556 Rowan Jen, Tim Jones, Gregg Kammerer, Ben Katz, Liyiming Ke, Mairbek Khadikov, Chandra
557 Kuchi, Marinda Lamb, Devin LeBlanc, Brendon LeCount, Sergey Levine, Xinyu Li, Adrian Li-
558 Bell, Vladislav Lialin, Zhonglin Liang, Wallace Lim, Yao Lu, Enyu Luo, Vishnu Mano, Nandan
559 Marwaha, Aikys Mongush, Liam Murphy, Suraj Nair, Tyler Patterson, Karl Pertsch, Allen Z. Ren,
560 Gavin Schelske, Charvi Sharma, Baifeng Shi, Lucy Xiaoyang Shi, Laura Smith, Jost Tobias Sprin-
561 genberg, Kyle Stachowicz, Will Stoeckle, Jiaming Tang, Jimmy Tanner, Shalom Tekeste, Marcel
562 Torne, Kyle Vedder, Quan Vuong, Anna Walling, Haohuan Wang, Jason Wang, XuDong Wang,
563 Chris Whalen, Samuel Whitmore, Blake Williams, Charles Xu, Sukwon Yoo, Lili Yu, Wuming
564 Zhang, Zhuoyang Zhang, and Ury Zhilinsky. $\pi_{0.7}$: a steerable generalist robotic foundation model
565 with emergent capabilities, 2026. URL <https://arxiv.org/abs/2604.15483>.
- 566 Huiwon Jang, Sihyun Yu, Heeseung Kwon, Hojin Jeon, Younggyo Seo, and Jinwoo Shin. Con-
567 textvla: Vision-language-action model with amortized multi-frame context, 2025. URL <https://arxiv.org/abs/2510.04246>.
568
- 569 Li Ji, Siyin Wang, Pengfang Qian, Xiaopeng Yu, Yihai Tian, Zhaoye Fei, Jingjing Gong, and Xipeng
570 Qiu. Hime: Hierarchical embodied memory for long-horizon vision-language-action control,
571 2026. URL <https://arxiv.org/abs/2607.03449>.
572
- 573 Yuanchen Ju, Yongyuan Liang, Yen-Jen Wang, Nandiraju Gireesh, Yuanliang Ju, Seungjae Lee,
574 Qiao Gu, Elvis Hsieh, Furong Huang, and Koushil Sreenath. Momagraph: State-aware unified
575 scene graphs with vision-language model for embodied task planning, 2026. URL <https://arxiv.org/abs/2512.16909>.
576
- 577 Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair,
578 Rafael Rafailov, Ethan Foster, Grace Lam, Pannag Sanketi, Quan Vuong, Thomas Kollar, Ben-
579 jamin Burchfiel, Russ Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang, and Chelsea Finn.
580 Openvla: An open-source vision-language-action model, 2024. URL <https://arxiv.org/abs/2406.09246>.
581
- 582 Huashuo Lei, Wenxuan Song, Huarui Zhang, Jieyuan Pei, Jiayi Chen, Haodong Yan, Han Zhao,
583 Pengxiang Ding, Zhipeng Zhang, Lida Huang, Donglin Wang, Yan Wang, and Haoang Li.
584 Robomemarena: A comprehensive and challenging robotic memory benchmark, 2026. URL
585 <https://arxiv.org/abs/2605.10921>.
586
- 587 Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and
588 Andy Zeng. Code as policies: Language model programs for embodied control, 2023. URL
589 <https://arxiv.org/abs/2209.07753>.
590
- 591 Rutav Shah, Yisu Li, Femi Bello, Yuke Zhu, and Roberto Martín-Martín. Memory retrieval in
592 visuomotor policies for long-horizon robot control, 2026. URL <https://arxiv.org/abs/2606.25136>.
593

Hao Shi, Bin Xie, Yingfei Liu, Lin Sun, Fengrong Liu, Tiancai Wang, Erjin Zhou, Haoqiang Fan, Xiangyu Zhang, and Gao Huang. Memoryvla: Perceptual-cognitive memory in vision-language-action models for robotic manipulation, 2026. URL <https://arxiv.org/abs/2508.19236>.

Lucy Xiaoyang Shi, Brian Ichter, Michael Equi, Liyiming Ke, Karl Pertsch, Quan Vuong, James Tanner, Anna Walling, Haohuan Wang, Niccolo Fusai, Adrian Li-Bell, Danny Driess, Lachy Groom, Sergey Levine, and Chelsea Finn. Hi robot: Open-ended instruction following with hierarchical vision-language-action models, 2025. URL <https://arxiv.org/abs/2502.19417>.

Ajay Sridhar, Jennifer Pan, Satvik Sharma, and Chelsea Finn. Memer: Scaling up memory for robot control via experience retrieval, 2025. URL <https://arxiv.org/abs/2510.20328>.

Xiaomi Robotics Team, Jun Guo, Piaopiao Jin, Jason Li, Peiyan Li, Yingyan Li, Futeng Liu, Wanli Peng, Optimus Qin, Yifei Su, Nan Sun, Qiao Sun, Runze Suo, Heyun Wang, Yunhong Wang, Rujie Wu, Caoyu Xia, Lina Zhang, Jack Zhao, Guoliang Chen, Wenlong Chen, Xinze He, Bin Li, Qing Li, Zhuorong Li, Heng Qu, Wenxuan Song, Diyun Xiang, Yifan Xie, Peiran Xu, Hangjun Ye, Wen Ye, Han Zhao, and Quanyun Zhou. Xiaomi-robotics-1: Scaling vision-language-action models with over 100k hours of real-world trajectories, 2026. URL <https://arxiv.org/abs/2607.15330>.

Marcel Torne, Karl Pertsch, Homer Walke, Kyle Vedder, Suraj Nair, Brian Ichter, Allen Z. Ren, Haohuan Wang, Jiaming Tang, Kyle Stachowicz, Karan Dhabalia, Michael Equi, Quan Vuong, Jost Tobias Springenberg, Sergey Levine, Chelsea Finn, and Danny Driess. Mem: Multi-scale embodied memory for vision language action models, 2026. URL <https://arxiv.org/abs/2603.03596>.

Ganlin Yang, Zhangzheng Tu, Yuqiang Yang, Sitong Mao, Junyi Dong, Tianxing Chen, Jiaqi Peng, Jing Xiong, Jiafei Cao, Jifeng Dai, Wengang Zhou, Yao Mu, and Tai Wang. Eventvla: Event-driven visual evidence memory for long-horizon vision-language-action policies, 2026. URL <https://arxiv.org/abs/2606.20092>.

Sizhe Zhao, Haozhe Xie, Weiyu Zhao, Chenchu Zhang, Huan Wang, Chenyang Wang, Qinglin Liu, and Shengping Zhang. Memory as plans: World-action modeling with memory-grounded planning, 2026. URL <https://arxiv.org/abs/2609.11561>.

A ALGORITHMS

This appendix provides the pseudocode for the two procedures described in the main text. Algorithm 1 summarizes one VLM query and the corresponding updates to the visual, event, and executable memories. Algorithm 2 details the self-labeling rollout and task-grounded verification used in Stage 2.

Algorithm 1: Streaming memory update at query t

Require: Previous memory $\mathcal{M}_{t-\Delta t}$, visual buffer $\mathcal{V}_{t-\Delta t}$, image chunk $\mathcal{C}_t$, instruction L

- 1 $\mathcal{V}_t \leftarrow \text{Append}(\mathcal{V}_{t-\Delta t}, \mathcal{C}_t)$;
 - 2 $(\Delta \mathcal{E}_t, \Delta \mathcal{P}_t, \ell_t) \leftarrow g_\theta(\mathcal{M}_{t-\Delta t}, \mathcal{C}_t, L)$;
 - 3 $\mathcal{P}_t \leftarrow \text{Append}(\mathcal{P}_{t-\Delta t}, \Delta \mathcal{P}_t)$;
 - 4 $\tilde{\mathcal{E}}_{t-\Delta t} \leftarrow \begin{cases} \emptyset, & \Delta \mathcal{P}_t \neq \emptyset, \\ \mathcal{E}_{t-\Delta t}, & \text{otherwise} \end{cases}$;
 - 5 $\mathcal{E}_t \leftarrow \text{Append}(\tilde{\mathcal{E}}_{t-\Delta t}, \Delta \mathcal{E}_t)$;
 - 6 **if** $\Delta \mathcal{E}_t \neq \emptyset$ **or** $\Delta \mathcal{P}_t \neq \emptyset$ **then**
 - 7 $\mathcal{V}_t \leftarrow \emptyset$;
 - 8 **return** $\mathcal{M}_t = (\mathcal{V}_t, \mathcal{E}_t, \mathcal{P}_t)$ and ℓ_t ;
-

Algorithm 2: Stage 2 self-labeling with task-grounded verification

Require: Stage 1 model g_{θ_1} , query interval Δt , maximum timestep T , training trajectories $\mathcal{T}$ with reference updates $\{(t_i, \Delta \mathcal{E}_i^{\text{ref}}, \Delta \mathcal{P}_i^{\text{ref}})\}_{i=1}^N$

Ensure : Verified self-labeled trajectories $\mathcal{D}_{\text{self}}$

```

1  $\mathcal{D}_{\text{self}} \leftarrow \emptyset$ ;
2 foreach  $\tau \in \mathcal{T}$  do
3   Initialize  $\mathcal{M}_0$  and temporary annotations  $\mathcal{D}_\tau \leftarrow \emptyset$ ;
4    $\text{valid} \leftarrow \text{True}$ ;  $t \leftarrow 0$ ;
5   while  $t < T$  and  $\text{valid}$  do
6      $t \leftarrow t + \Delta t$ ;
7      $(\Delta \mathcal{E}_t, \Delta \mathcal{P}_t, \ell_t) \leftarrow g_{\theta_1}(\mathcal{M}_{t-\Delta t}, \mathcal{C}_t, L)$ ;
8     Find  $i$  such that  $t \in [\frac{t_{i-1}+t_i}{2}, \frac{t_i+t_{i+1}}{2})$ ;
9     if  $\Delta \mathcal{E}_t \neq \Delta \mathcal{E}_i^{\text{ref}}$  or  $\Delta \mathcal{E}_t \in \mathcal{E}_{t-\Delta t}$  then
10       $\Delta \mathcal{E}_t \leftarrow \emptyset$ ;
11     if  $\Delta \mathcal{P}_t \neq \Delta \mathcal{P}_i^{\text{ref}}$  or  $\Delta \mathcal{P}_t \in \mathcal{P}_{t-\Delta t}$  then
12       $\Delta \mathcal{P}_t \leftarrow \emptyset$ ;
13     Update  $\mathcal{M}_t = (\mathcal{V}_t, \mathcal{E}_t, \mathcal{P}_t)$  with  $(\Delta \mathcal{E}_t, \Delta \mathcal{P}_t)$  as in Algorithm 1;
14     Add  $(\mathcal{M}_{t-\Delta t}, \mathcal{C}_t, L, \Delta \mathcal{E}_t, \Delta \mathcal{P}_t, \ell_t)$  to  $\mathcal{D}_\tau$ ;
15     foreach unchecked  $i$  with  $t \geq \frac{t_i+t_{i+1}}{2}$  do
16       if  $\Delta \mathcal{E}_i^{\text{ref}} \notin \mathcal{E}_t$  or  $\Delta \mathcal{P}_i^{\text{ref}} \notin \mathcal{P}_t$  then
17          $\text{valid} \leftarrow \text{False}$ ; break;
18       Mark  $i$  as checked;
19   if  $\text{valid}$  then
20      $\mathcal{D}_{\text{self}} \leftarrow \mathcal{D}_{\text{self}} \cup \{\mathcal{D}_\tau\}$ ;
21 return  $\mathcal{D}_{\text{self}}$ ;

```

B PAL TRAINING DETAILS**B.1 STAGE 1 TRAINING SETUP**

We build **PAL** upon MemER (Sridhar et al., 2025). Table 4 summarizes the network architecture and the Stage 1 training hyperparameters.

B.2 STAGE 2 SELF-LABELING

Stage 2 starts from the Stage 1 LoRA adapter and keeps the optimizer, batch size, and LoRA configuration above. It has two modes, selected according to how reliable the Stage 1 model already is.

Verifier. The Stage 2 verifier checks event and subtask-list updates independently against their reference updates. All non-coordinate text must match exactly (100% correct; paraphrases, omissions, and extra text are rejected). When coordinates are present, the prediction must include the corresponding coordinate fields, and each predicted-reference pair must have an L2 distance of strictly less than 10 pixels, i.e., $\|\mathbf{x}^{\text{pred}} - \mathbf{x}^{\text{ref}}\|_2 < 10$; without coordinates, the exact-text criterion alone applies.

We use one-pass self-labeling for the simulation experiments because the Stage 1 model already produces reliable rollouts: the first pass accepts most training episodes, so no further labeling rounds are needed. We use cyclic self-labeling for the real-world experiments because the first pass is less reliable; it retains only 31 of 200 episodes in the first round and improves to 170 in the final round. The change in acceptance over cycles is shown in Figure 5.

One-pass labeling. When the Stage 1 model is already strong enough, we label the training set once and then train. The Stage 1 model rolls out every training episode, and only episodes accepted

Table 4: Network architecture and Stage 1 training hyperparameters of PAL.

Configurations	Values
<i>Network Architecture</i>	
Base VLM	Qwen3-VL-4B-Instruct
Low-level policy	Pretrained MemER VLA
Action chunk length	20
Fine-tuning	LoRA on language-model linear layers
Frozen modules	Vision tower and aligner
LoRA rank / α / dropout	16 / 32 / 0.05
Trainable parameters	33.0M (0.74%)
<i>Training Hyperparameters</i>	
Hardware	8 $\times$ NVIDIA A800
Parallelism	DeepSpeed ZeRO-2
Optimizer	Fused AdamW
Epochs	8
Training time	$\sim$ 8.9 hours
Learning rate	8×10^{-5}
Schedule	Cosine, warmup ratio 0.05
Adam (β_1, β_2)	(0.9, 0.95)
Weight decay	0.1
Gradient clipping	1.0
Per-device batch size	4
Gradient accumulation	1
Global batch size	32
Precision	bfloat16
Attention	FlashAttention
Max sequence length	3200

by the verifier in Algorithm 2 are kept. The resulting corpus is the accepted self-labeled episodes together with the original supervised labels of the tasks that are not self-labeled (BinFill, ButtonUnmask, PickXtimes, StopCube, and SwingXtimes). Non-empty accepted updates are repeated three extra times, and PatternLock and RouteStick are upweighted by a factor of four. We then train for 2 epochs on 8 GPUs. In the reported run this corpus contains 178,855 rows and the training run takes about 5.5 hours. Of the 100 training episodes per self-labeled task, the verifier accepts 100 for MoveCube, 99 for PickHighlight and VideoUnmask, 98 for InsertPeg, PatternLock, and RouteStick, 92 for VideoPlaceButton, 90 for ButtonUnmaskSwap, 86 for VideoPlaceOrder and VideoUnmaskSwap, and 84 for VideoRepick. The first self-labeling pass already accepts most episodes of every task, so this single pass provides enough data for training and no further labeling rounds are needed.

Cyclic labeling. When the Stage 1 model is still weak, a single labeling pass accepts too few episodes, so we alternate training and relabeling. The reported run self-labels PatternLock and VideoPlaceOrder, both on their conditioning videos, and keeps the original supervised labels for PickXtimes and VideoUnmask. Both the event and the subtask list are verified. An initial labeling pass is run with the Stage 1 adapter, after which training proceeds for 3000 steps in blocks of 200. After each block, the current checkpoint relabels every episode of the two self-labeled tasks; accepted episodes replace their previous labels, while episodes that fail verification fall back to their original supervised labels. The next block resumes from that checkpoint and trains on the refreshed corpus. This run uses 15 such rounds. An *accepted* episode is one that passes the verifier of Sec. 4.3 and can therefore be used for training. A *normal rollout success* is an episode whose final list is already correct without the verifier. Acceptance rises quickly: the first round keeps 31 of 200 episodes, and the last round keeps 170. Non-empty accepted updates are repeated three extra times, as in the one-pass mode. The final corpus contains 40,144 rows, and the 15 training rounds take about 0.7 hours in total. Figure 5 shows how both quantities change over the 15 cycles.




Figure 5: Accepted episodes and normal rollout successes over the 15 self-labeling cycles. Each cycle relabels all 200 episodes of PatternLock and VideoPlaceOrder.

Table 5: Full RoboMME results. Success rates (%) are reported for all 16 tasks, grouped by how the executable subtask list is obtained. The highest score for each task is in bold.

Method	Short-video							Long-video		In-context				Episode-start			AVG
	Video Umk	Video Repick	Video UmkS	Pattern Lock	Route Stick	Move Cube	Insert Peg	Video PlcBtn	Video PlcOrd	Button Umk	Button UmkS	Pick HighL	Stop Cube	Pick Xtimes	Bin Fill	Swing Xtimes	
Baselines																	
$\pi_{0.5}$	20.44	0.44	18.67	2.89	4.67	26.00	1.56	31.11	25.78	22.22	6.67	11.33	6.67	42.89	30.00	35.56	17.93
EventVLA	66.00	30.00	17.60	4.00	27.60	44.00	4.00	22.00	16.00	60.00	12.00	12.00	46.00	48.00	6.00	10.00	26.57
FrameSamp-Modul	32.67	30.44	24.44	53.56	66.67	77.78	7.56	60.00	32.00	25.11	18.22	22.89	42.00	87.33	39.56	92.00	44.51
MemER	81.33	25.33	38.00	16.67	12.00	82.67	6.67	30.00	26.00	72.00	21.33	70.67	0.00	79.33	56.67	59.33	42.38
GPT-6-astra + VLA	78.00	32.00	28.00	24.00	4.00	54.00	2.00	50.00	82.00	32.00	0.00	74.00	6.00	88.00	38.00	44.00	39.75
Streaming Reasoning																	
Stage 1 Only	70.80	30.80	17.60	33.60	22.00	42.40	2.80	18.80	16.00	62.40	12.38	48.40	6.80	86.40	51.61	74.80	37.35
Sparse Label	76.40	28.40	18.80	24.40	16.80	81.20	11.20	40.88	20.00	74.91	19.24	58.00	2.40	87.60	56.40	77.60	43.39
PAL (ours)	95.20	32.00	35.20	50.80	22.40	76.40	4.80	80.00	83.20	90.80	20.40	68.40	5.20	87.20	65.20	79.60	56.05

C SIMULATED EXPERIMENTAL DETAILS

C.1 FULL SIMULATION RESULTS

The per-task results in Table 5 show that **PAL** performs best on VideoUnmask (95.20%), ButtonUnmask (90.80%), VideoPlaceButton (80.00%), and VideoPlaceOrder (83.20%). The margins over the strongest baseline are 13.87, 18.80, 20.00, and 39.20 percentage points, respectively. **PAL** is comparable to the strongest baseline on PatternLock (50.80% vs. 53.56%), MoveCube (76.40% vs. 82.67%), PickHighlight (68.40% vs. 70.67%), and PickXtimes (87.20% vs. 87.33%). It also leads BinFill with 65.20%, compared with 56.67% for MemER.

PAL remains weaker on RouteStick (22.40% vs. 66.67%) and SwingXtimes (79.60% vs. 92.00%). Its largest limitations are InsertPeg (4.80%) and StopCube (5.20%), where the best alternative scores are 11.20% and 46.00%, respectively. VideoRepick and the two swap tasks also remain difficult: the best scores are 32.00% on VideoRepick, 38.00% on VideoUnmaskSwap, and 21.33% on ButtonUnmaskSwap.

C.2 DATASET CONSTRUCTION

We construct sparse labels from each simulator episode. A sparse label is created only when the event memory or the executable subtask list should change. The way we construct the list depends on the task type. We use the simulator’s oracle event annotations only to determine what happened in the episode and when the corresponding information becomes available; the model is still trained to infer the event from visual observations.

Table 6: **PAL** inference and evaluation protocol on RoboMME.

Configurations	Values
Conditioning-video sampling	Every 4 timesteps
Execution action chunk	16
VLM query	Once per action chunk
Visual buffer input	4 uniformly sampled frames
Policy seeds	42, 43, 44, 45, 46

We use the following three cases:

1. **Episode-start tasks.** The ordered `subtask_list` can be determined at the beginning of the episode. We therefore construct the full list at the first query and keep it unchanged during the rest of the episode. These tasks do not need an intermediate event label for constructing the list.
2. **Tasks with changing environments.** We first read the oracle event sequence from the simulator. For each task, a task-specific rule-based procedure determines when an event can be inferred from the observations and which subtask or subtasks should be added to the `subtask_list`. Before enough evidence is available, the list is left unchanged. After the event can be inferred, the sparse label first records the event; when the information is sufficient to specify the next action, a later sparse label constructs or extends the `subtask_list`. The rules for this timing and mapping are defined separately for each task.
3. **StopCube.** This task does not require a separate event annotation. At each timestep, we set the executable `subtask_list` to the current subtask itself, so the list and the subtask are identical.

Thus, depending on the task, a sparse sample can supervise an event update, a `subtask_list` update, or both. Between two update points, the memory state is kept unchanged. Stage 1 subsequently expands these sparse update points into a timing band with multiple query timesteps, as described in Sec. 4; this expansion is separate from the sparse-label construction above.

C.3 IMPLEMENTATION DETAILS AND BASELINES

We first describe the **PAL** implementation and evaluation protocol, followed by the configurations of the comparison methods.

PAL. The **PAL** training setup is described in Appendix B. Table 6 summarizes the inference and evaluation protocol. We evaluate **PAL** with 50 episodes per task for each seed. For tasks with a conditioning video, we feed the video to the robot as a stream, even though the entire video is available at the beginning. This setting is practically meaningful because it evaluates the model’s ability to reason in real time.

Baselines. For the $\pi_{0.5}$, MemER, and FrameSamp-Modul baselines, we use the results reported by RoboMME (Dai et al., 2026); the original paper evaluates each method using three policy checkpoints and three seeds per checkpoint, yielding nine evaluation runs in total.

EventVLA. We adapt the original EventVLA implementation by using the frames annotated as event or `subtask_list` updates in **PAL** as keyframes. We also implement a streaming keyframe-selection procedure that selects keyframes from the history video as it is observed. All other settings follow the default EventVLA configuration. For long-horizon tasks, we increase the maximum number of stored keyframes to eight to accommodate tasks that require more historical evidence. The model is trained for 100k steps, with a maximum of eight keyframes at inference time.

GPT-6-*astra* + VLA. We adopt a structure similar to that of (Torne et al., 2026), in which memory is maintained as a language description of the current environment. At each query, the VLM reads the four most recent frames and decides the next subtask for the low-level VLA to execute. For

the conditioning video, we sample one frame every 33 frames, using at most 24 frames; for videos shorter than 100 frames, we provide at least three frames. This baseline uses the same low-level VLA as [PAL](#) but is evaluated with one seed because repeated API calls incur substantial evaluation time.

C.4 MORE ANALYSIS

C.4.1 GPT-6-ASTRA + VLA

We analyze the results of GPT-6-astra + VLA to identify its failure modes.

Failure mode 1: insensitivity to subtle environmental changes. The model cannot reliably detect small changes in the scene. For example, whether a button has been pressed is visually subtle, so the model often cannot judge whether the press succeeded and therefore does not output the next subtask. This is reflected in its failures on ButtonUnmask and ButtonUnmaskSwap, where it performs notably worse than on the similar tasks VideoUnmask and VideoUnmaskSwap, which do not require pressing a button.

Failure mode 2: loss of rapidly changing historical information. GPT-6-astra achieves the second-best performance on the Long-video tasks, yet fails on Short-video tasks such as RouteStick and PatternLock. These tasks are information-dense and their evidence spans only a short interval, so uniform downsampling does not retain enough information.

Failure mode 3: unstable high-level-low-level interface. The model does not interface reliably with the low-level policy. On tasks such as BinFill and VideoUnmask, the high-level VLM must specify which cube or box to pick. However, the bounding box it outputs jumps frequently from one location to another, which interrupts the motion of the low-level VLA.

C.4.2 [PAL](#)

[PAL](#) also performs poorly on InsertPeg, StopCube, and RouteStick. InsertPeg requires precise manipulation to align and insert the peg. In the 250 evaluation episodes, 101 ended in timeout because the peg was not inserted within the allowed horizon, contributing to the low success rate of 4.8%. StopCube requires the robot to anticipate the motion of a dynamically moving object and react quickly enough to press the button at the right moment. These fine-control and rapid-reaction requirements are outside the main scope of this work, which focuses on memory update and long-horizon reasoning. RouteStick exposes a different limitation: the high-level subtask selection and the low-level action execution are not always well coordinated. In an auxiliary experiment, reducing the action-chunk size from 16 to 8 increased the RouteStick success rate from 22.4% to 46.0%, suggesting that finer execution granularity substantially improves this coordination.

D REAL-WORLD EXPERIMENTAL DETAILS

D.1 DATA COLLECTION AND PROCESSING

We first use gripper actions and intervals in which the robot arm remains stationary to obtain coarse temporal annotations of each trajectory, providing initial estimates of action and subtask boundaries without frame-level manual annotation. We then provide these coarse annotations together with the recorded trajectory videos to GPT-5.6-Sol for fine-grained annotation; GPT-5.6-Sol refines the temporal boundaries and produces the task-level labels used to construct the event and executable subtask supervision. Finally, we remove the stationary portions identified during this process while preserving frames around gripper transitions, subtask boundaries, and task completion, synchronize the remaining front and wrist RGB observations with the robot states and actions, crop and resize both views using fixed camera-specific regions, and re-index the trajectory. Each retained action is rewritten to target the next retained observation, yielding the final action-aligned trajectories used in our experiments.

Table 7: Difficulty settings for the real-world tasks. PickXtimes uses three levels, while VideoUnmask has no difficulty distinction.

Task	Level 1	Level 2	Level 3	Level 4
PickXtimes	2 successful picks	3 successful picks	4 successful picks	–
VideoPlaceOrder	2 placements	3 placements	4 placements	5 placements
PatternLock	1 movement	2 movements	3 movements	4 movements
VideoUnmask	No difficulty levels			

D.2 SETTINGS AND IMPLEMENTATION DETAILS

Table 7 summarizes how task difficulty is varied in our real-world evaluation. Difficulty is controlled by the amount of task-relevant history that must be retained. In PickXtimes, the robot repeatedly picks up a green box and returns it to its original position, so the required number of successful picks must be remembered. In VideoPlaceOrder, the robot watches a video in which a green box is placed at several locations and then moves the box to the location specified by its order in the demonstration. In PatternLock, the robot watches a path over a 3×3 grid and then reproduces the same sequence of movements. In VideoUnmask, the robot watches three colored cubes being covered, remembers the color-position mapping, and then points to the requested cube; this task has no difficulty levels. The training data, evaluation trajectories, and task configurations are generated using GPT-6-astra. We collect 100 training episodes for each task and 200 episodes for VideoPlaceOrder because each VideoPlaceOrder episode is relatively short.

We use a Franka Research 3 robot arm equipped with a front-facing camera and a wrist-mounted camera. Inference runs on a single NVIDIA RTX 4090 GPU, and the inference computer communicates with the robot through an SSH-based ROS bridge. The robot is controlled by a low-level impedance controller that executes the VLA-generated commands at approximately 950 Hz. Excluding network latency, the VLA generates action commands at approximately 10–12 Hz.

D.3 EVALUATION PROTOCOL

We apply the same task-specific evaluation criteria to all methods to ensure fair comparisons. VideoPlaceOrder, PatternLock, and VideoUnmask use pre-recorded conditioning videos to reduce the manual effort required during evaluation. For PickXtimes, Levels 1–3 contain 5, 3, and 2 episodes, respectively. For VideoPlaceOrder and PatternLock, Levels 1–4 contain 3, 3, 2, and 2 episodes, respectively. VideoUnmask has no difficulty levels and uses 10 episodes under its fixed setting. Each episode is repeated three times, yielding 30 evaluations per task. Across all tasks, an episode is marked as a failure if it exceeds the limit of 1,300 timesteps or triggers the robot arm’s safety lock.

PickXtimes. An episode is successful if the robot lifts the green box off the table exactly the specified number of times and subsequently keeps its arm raised for 100 timesteps without reaching down to pick up the box again. Any additional contact with the box after the required number of repetitions is counted as a failure. To reduce evaluation time, a trial is also terminated and counted as a failure if any individual pick-and-place cycle takes more than 650 timesteps. The robot arm and the green box start from the same fixed positions in every episode.

VideoPlaceOrder. An episode is successful if the green box is on the designated target area when the gripper releases it. A placement is considered valid if the box’s vertical projection overlaps the target area. An episode fails if the box is outside all target areas upon release or touches an incorrect target area. The green box is initialized at the center of the workspace, and the robot arm starts from the same fixed pose in every episode.

PatternLock. An episode is successful if the gripper visits the grid points in the prescribed order. A grid point is considered visited when any part of the gripper passes directly above it. An episode fails if the gripper leaves the grid region or passes over an incorrect grid point. At the beginning of each episode, the gripper is closed and the robot arm is placed at the initial point of the target trajectory.




Figure 6: Real-world rollout visualizations of **PAL**. From top to bottom, the rows show PatternLock, PickXtimes, VideoPlaceOrder, and VideoUnmask, with eight frames sampled along each rollout.

VideoUnmask. An episode is successful if the robot points to the correct box after the demonstration video. An episode is failed if the robot points to the wrong box. The robot arm starts from the same fixed pose in every episode.

D.4 ROLLOUT VISUALIZATION

Fig. 6 shows representative real-world rollouts of **PAL**, with one task per row and eight frames sampled along each rollout. In PatternLock, the robot watches a path over the 3×3 grid and then reproduces the same sequence of movements. In PickXtimes, the robot repeatedly picks up the green box and returns it to its original position, stopping after the specified number of picks. In VideoPlaceOrder, the robot watches a video in which a green box is placed at several locations and then moves the box to the location specified by its order in the demonstration. In VideoUnmask, the robot watches three colored cubes being covered and then points to the requested cube.

D.5 MORE ANALYSIS

The baseline failures reveal distinct limitations in task-progress tracking and visual target grounding (Table 3). On PickXtimes, MemER achieves 0.00% success because it fails to track how many pick cycles have been completed and repeatedly outputs the initial subtask instead of advancing to the next one. On VideoUnmask, FrameSamp-Modul achieves only 30.00% success, close to the 33.33% expected from randomly choosing one of the three target boxes. In the observed rollouts, it repeatedly points to the middle box rather than grounding its choice in the requested cube.

E PROMPT

We show two inference queries. The first is from the video phase of VideoPlaceOrder, where the model updates the subtask list from the conditioning video. The second is from the execution phase of PatternLock, where the model selects the current subtask from the subtask list. In both, the system message fixes the output schema, the user message gives the task, the current memory, and the selected frames, and the output is the model’s response at this query. The four images of each query are supplied at the timesteps marked below and are shown in Fig. 7 and Fig. 8; their file paths are omitted.

```
System:
Update the robot subtask list from selected conditioning images and inferred pre-phase
events.
Report the completed event inferred from the images; do not predict the next event.

Return JSON with:
- subtask_list: ordered subtasks inferred so far
- previous_phase_event: newly completed event for the current timestep, or null when there
  is no new event; do not copy an inferred event

User:
Task goal: watch the video carefully, then place the red cube on the second target it was
previously placed on
Phase: video watching
Visual context: front-view camera; left/right are mirrored, so reason in robot-centered
coordinates.
Use image coordinates as [y, x] with origin at the top-left.
Previous subtask list, last updated unknown: none
Inferred pre-phase events at sampled memory time, last updated 2.90 seconds (87 timesteps)
ago:
1. pick up the cube at <bbox>
2. drop the cube onto target at <bbox>
3. pick up the cube at <bbox>
4. drop the cube onto target at <bbox>
Images are selected around event or subtask-list update times and ordered by timestep.
Question: Update the subtask list from the selected conditioning images.
Grounded object state is included with the selected images.
Selected conditioning images:
Image 1 at timestep 349: <image>
Image 2 at timestep 378: <image>
Image 3 at timestep 407: <image>
Image 4 at timestep 436: <image>

Output:
{"subtask_list": "1. pick up the cube at <ungrounded> 2. drop the cube onto target at
<bbox>", "previous_phase_event": "press the button at <bbox>"}
```

```
System:
You select the robot’s current executable subtask after conditioning is complete.
```





Figure 7: The four conditioning images used in the video-phase prompt above, in timestep order: 349, 378, 407, and 436.




Figure 8: The four execution images used in the execution-phase prompt above, in timestep order: 58, 71, 83, and 96.

```

Decision rule: choose current_subtask from the subtask list using selected execution
  images. Keep subtask_list unchanged unless the visual evidence clearly shows it is
  incomplete or wrong. When unchanged, return subtask_list as null instead of copying the
  previous list.

Return JSON using only the fields supervised for this sample; not every field is required:
- subtask_list (optional): ordered subtasks for the task
- current_subtask: executable action at the current timestep
- previous_phase_event (optional): newly completed task-specific event, or null when
  unchanged

User:
Task goal: watch the video carefully, then use the stick attached to the robot to retrace
  the same pattern
Phase: execution
Visual context: front-view camera; left/right are mirrored, so reason in robot-centered
  coordinates.
Use image coordinates as [y, x] with origin at the top-left.
Subtask list, last updated unknown: 1. move backward 2. move right 3. move forward 4. move
  backward-right
Previous query current_subtask: 3. move forward
Inferred pre-phase events at sampled memory time, last updated unknown:
  none
Selected execution images are sampled from the current visual window through the current
  timestep.
Question: Predict the current executable subtask.
Grounded object state is included with the selected images.
Selected execution images:
Image 1 at timestep 58: <image>
Image 2 at timestep 71: <image>
Image 3 at timestep 83: <image>
Image 4 at timestep 96: <image>

Output:
{"current_subtask": "4. move backward-right"}

```